

## **Snowflake Learning Journey – Comprehensive Summary**

Date: December 8, 2025

Learning Path: Data Engineering Foundation with Snowflake

---

### **Table of Contents**

1. Snowflake Architecture
  2. Virtual Warehouses
  3. User & Role Management (RBAC)
  4. Connecting to Snowflake
  5. Data Types & Table Design
  6. Hands-on Practices
  7. Next Steps
- 

### **1. Snowflake Architecture**

#### **Overview**

Snowflake is a cloud-native data platform built on three core layers:

- Storage Layer – Scalable cloud object storage (S3, Azure Blob, GCS) storing data in Parquet format
- Compute Layer – Virtual warehouses (clusters) that execute queries independently
- Cloud Services Layer – Metadata management, authentication, query optimization, and transaction control

#### **Key Characteristics**

- Shared-Nothing Architecture – Storage and compute are decoupled; scale independently
- Snowflake Metadata Store – Centralized metadata repository tracking objects, privileges, and versioning
- Multi-Cloud Support – AWS, Azure, GCP
- Zero-Copy Clone – Instantly copy databases/schemas without duplicating data

#### **Advantages**

- Pay only for compute used (auto-suspend, auto-resume)
- Instant elasticity – scale up/down on demand
- Time-travel and fail-safe for data recovery
- Secure data sharing without copying data

---

## 2. Virtual Warehouses

### What is a Virtual Warehouse?

A virtual warehouse is a compute cluster that executes SQL queries. Each warehouse is independent and can run in parallel.

### Warehouse Types

#### Standard Warehouse

- Traditional multi-purpose compute
- Suitable for general SQL queries, ETL, analytics
- Example: 1–16 credits per hour

#### Snowpark-Optimized Warehouse

- Optimized for Snowpark (Python/Java/Scala) workloads
- Higher memory per compute unit
- Best for ML and data science jobs with large in-memory operations

#### Multi-Cluster Warehouse

- Auto-scales based on query queue demand
- Multiple clusters share the warehouse name
- Ideal for high-concurrency environments (many concurrent users/jobs)
- Example: Min 2, Max 10 clusters

### Creating & Monitoring Warehouses

Create a warehouse:

sql

**USE ROLE SYSADMIN;**

**CREATE OR REPLACE WAREHOUSE MY\_WH**

**WAREHOUSE\_SIZE = MEDIUM**

**AUTO\_SUSPEND = 60**

**AUTO\_RESUME = TRUE**

**INITIALLY\_SUSPENDED = FALSE;**

Monitor warehouse:

sql

```
SELECT * FROM SNOWFLAKE.ACCOUNT_USAGE.WAREHOUSE_METERING_HISTORY;
```

```
SELECT * FROM SNOWFLAKE.ACCOUNT_USAGE.QUERY_HISTORY;
```

Key parameters:

- WAREHOUSE\_SIZE – XS, S, M, L, XL, 2XL, 3XL, 4XL (credits per minute)
  - AUTO\_SUSPEND – Minutes before auto-pause (default 600)
  - AUTO\_RESUME – Automatically resume on query submission
  - MAX\_CLUSTER\_COUNT – For multi-cluster warehouses
- 

### **3. User & Role Management (RBAC)**

#### **Access Control Framework**

Snowflake uses Discretionary Access Control (DAC), Role-Based Access Control (RBAC), and User-Based Access Control (UBAC):

- DAC – Object owner can grant privileges
- RBAC – Privileges assigned to roles; roles assigned to users (primary method)
- UBAC – Direct user privileges (exceptions only)

#### **System-Defined Roles**

Role	Purpose	Can Do
ACCOUNTADMIN	Top-level admin	All privileges; account management
SECURITYADMIN	Security & grants	Grant/revoke privileges; manage users/roles
USERADMIN	User/role admin	Create users and roles; user lifecycle
SYSADMIN	System admin	Create databases, warehouses, schemas, tables
PUBLIC	Default role	Auto-assigned to all users
GLOBALORGADMIN	Org-level admin	Manage organization across accounts

## **Creating Custom Roles**

sql

```
USE ROLE SECURITYADMIN;
```

```
CREATE ROLE ROLE_APP_DB_OWNER COMMENT = 'Owns APP_DB and can grant privileges';
```

```
CREATE ROLE ROLE_APP_RW COMMENT = 'Read/write access to APP_DB';
```

```
CREATE ROLE ROLE_APP_RO COMMENT = 'Read-only access to APP_DB';
```

-- Attach into hierarchy for management

```
GRANT ROLE ROLE_APP_RW TO ROLE SYSADMIN;
```

```
GRANT ROLE ROLE_APP_RO TO ROLE SYSADMIN;
```

## **Creating Users & Assigning Roles**

sql

```
USE ROLE USERADMIN;
```

```
CREATE USER APP_DE_USER
```

```
PASSWORD = 'TempP@ssw0rd'
```

```
DEFAULT_ROLE = ROLE_APP_RW
```

```
MUST_CHANGE_PASSWORD = TRUE;
```

```
USE ROLE SECURITYADMIN;
```

```
GRANT ROLE ROLE_APP_RW TO USER APP_DE_USER;
```

## **Granting Privileges**

Grant on database:

sql

```
USE ROLE ROLE_APP_DB_OWNER;
```

```
GRANT USAGE ON DATABASE APP_DB TO ROLE ROLE_APP_RW;
```

```
GRANT USAGE ON SCHEMA APP_DB.APP_SCHEMA TO ROLE ROLE_APP_RW;
```

```
GRANT SELECT, INSERT, UPDATE, DELETE ON ALL TABLES IN SCHEMA APP_DB.APP_SCHEMA TO
ROLE ROLE_APP_RW;
```

Grant on future objects:

```
sql
```

```
GRANT SELECT ON FUTURE TABLES IN SCHEMA APP_DB.APP_SCHEMA TO ROLE ROLE_APP_RO;
```

### **Key Principles**

- USAGE privilege required to *see/reference* database or schema
  - SELECT/INSERT/UPDATE/DELETE required for table operations
  - Object owner can grant privileges on their objects
  - SECURITYADMIN (with MANAGE GRANTS) can grant any privilege globally
  - No privilege = no access (deny by default)
- 

## **4. Connecting to Snowflake**

### **Snowsight UI**

- Web-based IDE for writing and executing SQL
- Visual query results, charts, and dashboards
- Role switching and warehouse selection
- Built-in documentation
- Access: <https://<account>.snowflakecomputing.com>

### **SnowSQL**

Command-line interface for Snowflake.

Install & connect:

```
bash
```

```
# Install SnowSQL (macOS, Linux, Windows available)
```

```
snowsql -a <account_identifier> -u <username>
```

```
# Example commands
```

```
snowsql -a xy12345.us-east-1 -u sandeep_reporter
```

```
SELECT * FROM my_table;
```

```
!exit
```

### **VS Code Extension**

## Snowflake SQL Tools for VS Code

- Write SQL in IDE
- Execute queries
- Browse databases/schemas/tables
- Connect with credentials or key-pair auth

Setup:

1. Install Snowflake SQL Tools extension
2. Set connection in settings or .vscode/settings.json
3. Select warehouse/database/schema
4. Write and execute SQL

## Python Connector

Connect Python to Snowflake for ETL, data science, and automation.

python

```
from snowflake.connector import connect
```

```
conn = connect(  
    user='sandeep_reporter',  
    password='password',  
    account='xy12345.us-east-1',  
    warehouse='MY_WH',  
    database='APP_DB',  
    schema='APP_SCHEMA'  
)
```

```
cursor = conn.cursor()  
  
cursor.execute('SELECT * FROM CUSTOMERS LIMIT 10;')  
  
results = cursor.fetchall()  
  
for row in results:  
  
    print(row)  
  
cursor.close()
```

```
conn.close()
```

### JDBC Connection

For Java applications.

```
java
```

```
import java.sql.*;
```

```
public class SnowflakeConnection {
```

```
    public static void main(String[] args) throws Exception {
```

```
        Class.forName("net.snowflake.client.jdbc.SnowflakeDriver");
```

```
        String connectionString = "jdbc:snowflake://<account>.snowflakecomputing.com/" +
```

```
            "?warehouse=MY_WH&database=APP_DB&schema=APP_SCHEMA";
```

```
        Connection conn = DriverManager.getConnection(
```

```
            connectionString,
```

```
            "sandeep_reporter",
```

```
            "password"
```

```
);
```

```
        Statement stmt = conn.createStatement();
```

```
        ResultSet rs = stmt.executeQuery("SELECT * FROM CUSTOMERS;");
```

```
        while (rs.next()) {
```

```
            System.out.println(rs.getString(1));
```

```
        }
```

```
        rs.close();
```

```
        stmt.close();
```

```
        conn.close();
```

```
}
```

```
}
```

---

## 5. Data Types & Table Design

### Data Type Categories

Snowflake supports diverse data types grouped as:

#### Numeric

- NUMBER(precision, scale) – Default (38, 0)
- INT, BIGINT, SMALLINT, TINYINT
- FLOAT, DOUBLE, REAL
- DECIMAL, NUMERIC (synonyms for NUMBER)

#### String & Binary

- VARCHAR(length) – Default 16777216 bytes
- STRING, TEXT (synonyms for VARCHAR)
- CHAR(length) – Fixed length
- BINARY, VARBINARY

#### Date & Time

- DATE – Date only
- TIME – Time only
- TIMESTAMP\_NTZ – Timestamp without time zone
- TIMESTAMP\_LTZ – Timestamp with local time zone (stored as UTC)
- TIMESTAMP\_TZ – Timestamp with time zone (stored with TZ info)
- DATETIME (synonym for TIMESTAMP\_NTZ)

#### Boolean

- BOOLEAN – TRUE/FALSE

#### Semi-Structured

- VARIANT – JSON-like data (any JSON value)
- OBJECT – Key-value pairs
- ARRAY – List of values

#### Geospatial & Vector

- GEOGRAPHY – Points, lines, polygons (earth coordinates)
- GEOMETRY – Geometric objects (2D/3D)
- VECTOR(type, dimension) – ML embeddings (e.g., VECTOR(FLOAT, 128))

### **Example: CREATE TABLE with Various Data Types**

sql

```
CREATE OR REPLACE TABLE customer_orders (
    order_id      NUMBER(38,0)      PRIMARY KEY,
    customer_id   INT              NOT NULL,
    order_amount  NUMBER(10,2),
    discount_pct  FLOAT,
    email         VARCHAR(255),
    is_first_order BOOLEAN,
    order_date    DATE,
    order_time    TIME,
    created_at    TIMESTAMP_NTZ DEFAULT CURRENT_TIMESTAMP(),
    updated_at    TIMESTAMP_LTZ,
    notes         STRING,
    receipt_file  BINARY
);
```

### **Example: Semi-Structured Data**

sql

```
CREATE OR REPLACE TABLE click_events (
    event_id      NUMBER  PRIMARY KEY,
    user_id       VARCHAR,
    event_time    TIMESTAMP_TZ,
    event_data    VARIANT      -- Full JSON payload
);
```

```
INSERT INTO click_events VALUES (
    1,
    'user_123',
    CURRENT_TIMESTAMP(),
    PARSE_JSON('{"page": "home", "action": "click", "target": "signup_button"}')
);
```

```
-- Query VARIANT fields  
SELECT event_data:page, event_data:action FROM click_events;
```

#### **Example: Advanced Types**

sql

```
CREATE OR REPLACE TABLE advanced_data (  
    id      NUMBER PRIMARY KEY,  
    attributes OBJECT,           -- Key/value map  
    tags     ARRAY,             -- List  
    location GEOGRAPHY,        -- Point/polygon  
    embedding VECTOR(FLOAT, 128)   -- 128-dim vector  
);
```

---

## 6. Hands-on Practices

### **Practice 1: Create Database & Role Structure**

sql

```
USE ROLE SYSADMIN;  
CREATE DATABASE TEST_DB;  
CREATE SCHEMA TEST_DB.TEST_SCHEMA;
```

```
USE ROLE SECURITYADMIN;  
CREATE ROLE ROLE_DB_OWNER;  
GRANT ROLE ROLE_DB_OWNER TO ROLE SYSADMIN;
```

```
USE ROLE ROLE_DB_OWNER;  
CREATE DATABASE APP_DB;  
CREATE SCHEMA APP_DB.APP_SCHEMA;
```

### **Practice 2: Grant Privileges**

sql

```
USE ROLE SECURITYADMIN;
```

-- Create read-only role

```
CREATE ROLE DATA_READ_ONLY;
```

-- Grant USAGE on database and schemas

```
GRANT USAGE ON DATABASE APP_DB TO ROLE DATA_READ_ONLY;
```

```
GRANT USAGE ON SCHEMA APP_DB.APP_SCHEMA TO ROLE DATA_READ_ONLY;
```

-- Grant SELECT on tables

```
GRANT SELECT ON ALL TABLES IN SCHEMA APP_DB.APP_SCHEMA TO ROLE DATA_READ_ONLY;
```

```
GRANT SELECT ON FUTURE TABLES IN SCHEMA APP_DB.APP_SCHEMA TO ROLE DATA_READ_ONLY;
```

### Practice 3: Create User & Connect

sql

```
USE ROLE USERADMIN;
```

```
CREATE USER SANDEEP_REPORTER
```

```
PASSWORD = 'TempP@ssw0rd'
```

```
DEFAULT_ROLE = DATA_READ_ONLY
```

```
MUST_CHANGE_PASSWORD = TRUE;
```

```
USE ROLE SECURITYADMIN;
```

```
GRANT ROLE DATA_READ_ONLY TO USER SANDEEP_REPORTER;
```

Then connect via Snowsight, SnowSQL, VS Code, or Python as that user.

### Practice 4: Create & Monitor Warehouse

sql

```
USE ROLE SYSADMIN;
```

```
CREATE WAREHOUSE ANALYTICS_WH
```

```
WAREHOUSE_SIZE = MEDIUM
```

```
AUTO_SUSPEND = 60
```

```
AUTO_RESUME = TRUE;
```

-- Monitor

```
SELECT * FROM SNOWFLAKE.ACCOUNT_USAGE.WAREHOUSE_METERING_HISTORY
WHERE WAREHOUSE_NAME = 'ANALYTICS_WH'
ORDER BY START_TIME DESC;
```

#### **Practice 5: Load Data (CSV → Snowflake)**

sql

-- *Create file format*

```
CREATE FILE FORMAT FF_CSV
```

```
TYPE = 'CSV'
```

```
FIELD_DELIMITER = ','
```

```
SKIP_HEADER = 1;
```

-- *Create internal stage*

```
CREATE STAGE STG_DATA FILE_FORMAT = FF_CSV;
```

-- *Upload file (via SnowSQL or UI)*

```
-- PUT file:///path/to/data.csv @STG_DATA;
```

-- *Load into table*

```
COPY INTO CUSTOMERS
```

```
FROM @STG_DATA
```

```
FILES = ('data.csv')
```

```
ON_ERROR = 'ABORT_STATEMENT';
```

---

## **7. Next Steps & Learning Path**

### **Immediate Next Topics**

#### **1. Data Loading & Unloading**

- COPY INTO from internal/external stages
- Snowpipe for continuous loading
- UNLOAD for exporting data

#### **2. Transformations & Queries**

- Window functions (ROW\_NUMBER, RANK, LAG, LEAD)

- CTEs and nested queries
- Aggregations and GROUP BY
- Joins (INNER, LEFT, RIGHT, FULL OUTER, CROSS)

### 3. Advanced Features

- Snowpipe (streaming ingestion)
- Streams (change data capture)
- Tasks (scheduled SQL jobs)
- Dynamic Data Masking (DDM) for security

### 4. Performance Optimization

- Query profiling and EXPLAIN PLAN
- Clustering keys
- Micro-partitions and pruning
- Query optimization best practices

### 5. ETL/ELT Pipelines

- End-to-end pipeline design
- Error handling & logging
- Orchestration (Airflow, Prefect, dbt)
- Testing & monitoring

### 6. Data Sharing & Governance

- Secure Data Sharing
- Iceberg tables and Delta format
- Data catalog and lineage
- Audit logging

---

## Quick Reference: Common SQL Commands

### Database & Schema Management

sql

**CREATE DATABASE** db\_name;

**CREATE SCHEMA** db\_name.schema\_name;

**SHOW DATABASES;**

**SHOW SCHEMAS IN DATABASE** db\_name;

```
USE DATABASE db_name;
```

```
USE SCHEMA schema_name;
```

## Role & User Management

sql

```
CREATE ROLE role_name;
```

```
CREATE USER user_name PASSWORD = 'pwd';
```

```
GRANT ROLE role_name TO USER user_name;
```

```
GRANT USAGE ON DATABASE db_name TO ROLE role_name;
```

```
SHOW GRANTS TO ROLE role_name;
```

```
SHOW GRANTS ON DATABASE db_name;
```

## Warehouse Management

sql

```
CREATE WAREHOUSE wh_name WAREHOUSE_SIZE = MEDIUM;
```

```
ALTER WAREHOUSE wh_name SET WAREHOUSE_SIZE = LARGE;
```

```
SUSPEND WAREHOUSE wh_name;
```

```
RESUME WAREHOUSE wh_name;
```

```
SHOW WAREHOUSES;
```

## Table Operations

sql

```
CREATE TABLE table_name (col1 TYPE, col2 TYPE, ...);
```

```
INSERT INTO table_name VALUES (...);
```

```
SELECT * FROM table_name;
```

```
UPDATE table_name SET col1 = value WHERE condition;
```

```
DELETE FROM table_name WHERE condition;
```

```
DROP TABLE table_name;
```

## Query Insights

sql

-- *Query history*

```
SELECT * FROM SNOWFLAKE.ACCOUNT_USAGE.QUERY_HISTORY;
```

-- *Warehouse usage*

```
SELECT * FROM SNOWFLAKE.ACCOUNT_USAGE.WAREHOUSE_METERING_HISTORY;
```

-- Database usage

```
SELECT * FROM SNOWFLAKE.ACCOUNT_USAGE.DATABASE_USAGE_HISTORY;
```

---

## Resources

- Official Docs: <https://docs.snowflake.com>
  - Snowflake University: <https://learn.snowflake.com>
  - Community: <https://community.snowflake.com>
  - GitHub: <https://github.com/Snowflake-Labs>
- 

## Learning Summary

### Covered:

- Snowflake architecture and multi-cloud deployment
- Virtual warehouses (types, scaling, monitoring)
- User/role management and RBAC best practices
- Multiple connection methods (UI, CLI, Python, JDBC)
- Data types and table design
- Hands-on practices with real SQL commands

### Ready For:

- Building real ETL pipelines
  - Loading data from external sources
  - Designing role hierarchies for production
  - Performance tuning and query optimization
  - Advanced data engineering with Streams, Tasks, Snowpipe
- 

Last Updated: December 8, 2025

Learning Status: Foundation Complete | Ready for Advanced Topics